

ecg-solution

September 15, 2024

1 Imports

```
[ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
from sklearn.model_selection import cross_val_score
from imblearn.under_sampling import RandomUnderSampler
from sklearn.ensemble import RandomForestClassifier, \
    GradientBoostingClassifier, AdaBoostClassifier
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.model_selection import KFold
```

2 Exploración de archivos

```
[ ]: df = pd.read_csv('sample_data/mitbih_test.csv', header=None)

counts = {i: 0 for i in range(5)}


for _, row in df.iterrows():
    value = int(row[187])
    if 0 <= value <= 5:
        counts[value] += 1

for value, count in counts.items():
    print(f"Clase {value}: {count} ocurrencias")
```

```
counts_series = pd.Series(counts)

plt.figure(figsize=(8, 6))
counts_series.plot(kind='bar', color='red')
plt.title('Distribución de clases en el conjunto de datos')
plt.xlabel('Clases')
plt.ylabel('Número de instancias')
plt.show()
```

Clase 0: 18118 ocurrencias
 Clase 1: 556 ocurrencias
 Clase 2: 1448 ocurrencias
 Clase 3: 162 ocurrencias
 Clase 4: 1608 ocurrencias



```
[ ]: df = pd.read_csv('sample_data/mitbih_train.csv', header=None)

counts = {i: 0 for i in range(5)}
```

```

for _, row in df.iterrows():
    value = int(row[187])
    if 0 <= value <= 5:
        counts[value] += 1

for value, count in counts.items():
    print(f"Clase {value}: {count} ocurrencias")

counts_series = pd.Series(counts)


plt.figure(figsize=(8, 6))
counts_series.plot(kind='bar', color='red')
plt.title('Distribución de clases en el conjunto de datos')
plt.xlabel('Clases')
plt.ylabel('Número de instancias')
plt.show()

```

```

Clase 0: 72471 ocurrencias
Clase 1: 2223 ocurrencias
Clase 2: 5788 ocurrencias
Clase 3: 641 ocurrencias
Clase 4: 6431 ocurrencias

```



```
[ ]: df = pd.read_csv('sample_data/ptbdb_normal.csv', header=None)

counts = {i: 0 for i in range(5)}

for _, row in df.iterrows():
    value = int(row[187])
    if 0 <= value <= 5:
        counts[value] += 1

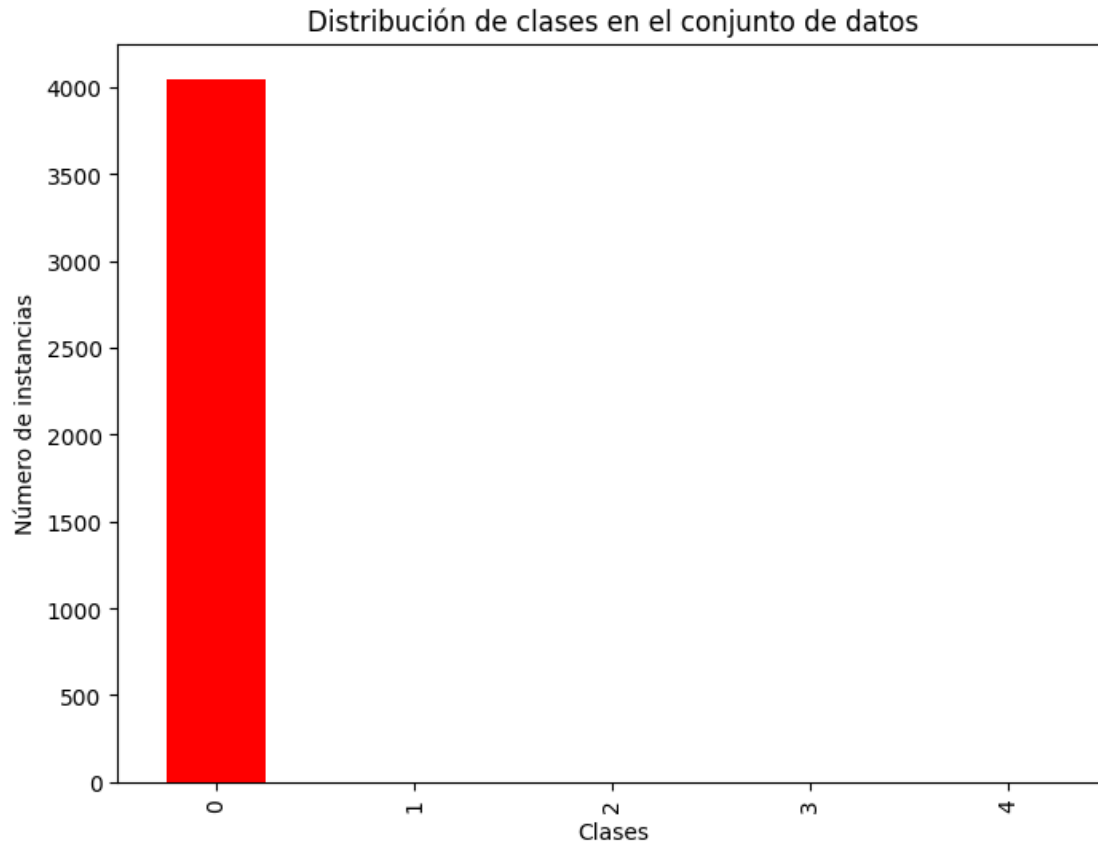
for value, count in counts.items():
    print(f"Clase {value}: {count} ocurrencias")

counts_series = pd.Series(counts)

plt.figure(figsize=(8, 6))
counts_series.plot(kind='bar', color='red')
plt.title('Distribución de clases en el conjunto de datos')
```

```
plt.xlabel('Clases')
plt.ylabel('Número de instancias')
plt.show()
```

Clase 0: 4046 ocurrencias
Clase 1: 0 ocurrencias
Clase 2: 0 ocurrencias
Clase 3: 0 ocurrencias
Clase 4: 0 ocurrencias



```
[ ]: df = pd.read_csv('sample_data/ptbdb_abnormal.csv', header=None)

counts = {i: 0 for i in range(5)}

for _, row in df.iterrows():
    value = int(row[187])
    if 0 <= value <= 5:
        counts[value] += 1
```

```


for value, count in counts.items():
    print(f"Clase {value}: {count} ocurrencias")

counts_series = pd.Series(counts)

plt.figure(figsize=(8, 6))
counts_series.plot(kind='bar', color='red')
plt.title('Distribución de clases en el conjunto de datos')
plt.xlabel('Clases')
plt.ylabel('Número de instancias')
plt.show()

```



Clase 0: 0 ocurrencias
 Clase 1: 10506 ocurrencias
 Clase 2: 0 ocurrencias
 Clase 3: 0 ocurrencias
 Clase 4: 0 ocurrencias


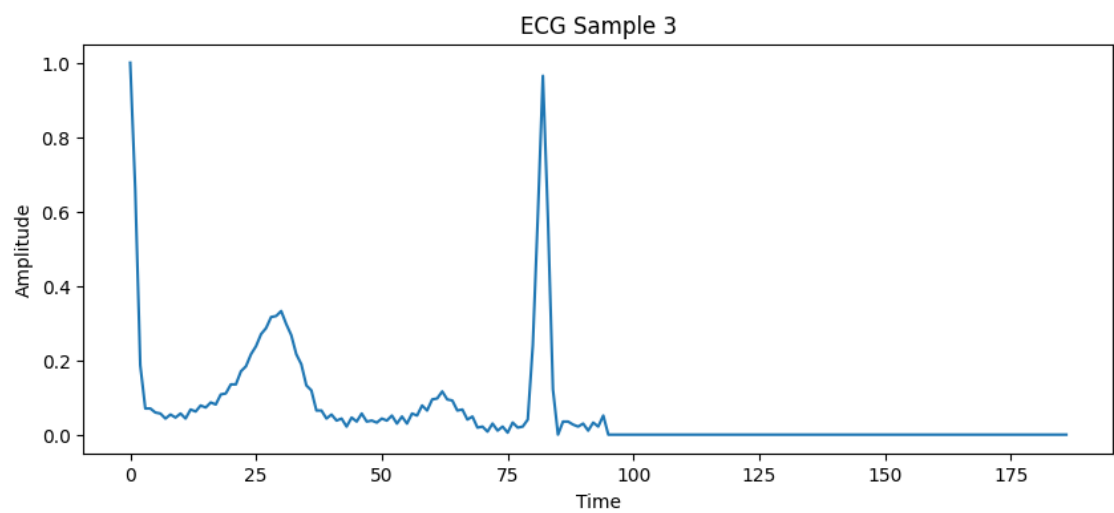



3 Visualización

```
[ ]: df = pd.read_csv('sample_data/mitbih_train.csv', header=None)

for i in range(5):
    plt.figure(figsize=(10, 4))
    plt.plot(df.iloc[i, :-1])
    plt.title(f'ECG Sample {i+1}')
    plt.xlabel('Time')
    plt.ylabel('Amplitude')
    plt.show()
```







4 Limpieza, Verificación de normalización y datos nulos

```
[ ]: df = pd.read_csv('sample_data/mitbih_test.csv', header=None)

df.isnull().sum()
print(df.describe())
```

| | 0 | 1 | 2 | 3 | 4 \ |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 21892.000000 | 21892.000000 | 21892.000000 | 21892.000000 | 21892.000000 |
| mean | 0.894410 | 0.761902 | 0.426627 | 0.221596 | 0.201676 |
| std | 0.234560 | 0.218659 | 0.228572 | 0.208711 | 0.177727 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.924260 | 0.683366 | 0.251197 | 0.050505 | 0.082873 |
| 50% | 0.990431 | 0.828996 | 0.432777 | 0.167630 | 0.147642 |
| 75% | 1.000000 | 0.912319 | 0.583991 | 0.347092 | 0.259211 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |

| | 5 | 6 | 7 | 8 | 9 \ |
|-------|--------------|--------------|--------------|--------------|--------------|
| count | 21892.000000 | 21892.000000 | 21892.000000 | 21892.000000 | 21892.000000 |
| mean | 0.209891 | 0.204805 | 0.200992 | 0.197634 | 0.196022 |
| std | 0.172194 | 0.177946 | 0.176142 | 0.170228 | 0.166707 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.087912 | 0.072663 | 0.065997 | 0.064516 | 0.068493 |
| 50% | 0.158111 | 0.144068 | 0.144509 | 0.150422 | 0.149029 |
| 75% | 0.287356 | 0.298453 | 0.294563 | 0.289907 | 0.282956 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 0.991429 |

| | ... | 178 | 179 | 180 | 181 \ |
|-------|-----|--------------|--------------|--------------|--------------|
| count | ... | 21892.000000 | 21892.000000 | 21892.000000 | 21892.000000 |

| | | | | | |
|------|-----|----------|----------|----------|----------|
| mean | ... | 0.004588 | 0.004327 | 0.004020 | 0.003789 |
| std | ... | 0.043128 | 0.042187 | 0.040255 | 0.039397 |
| min | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 50% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 75% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| max | ... | 0.980392 | 1.000000 | 0.966102 | 1.000000 |

| | | | | | | |
|-------|--------------|--------------|--------------|--------------|--------------|-------|
| | | 182 | 183 | 184 | 185 | 186 \ |
| count | 21892.000000 | 21892.000000 | 21892.000000 | 21892.000000 | 21892.000000 | |
| mean | 0.003638 | 0.003459 | 0.003166 | 0.003000 | 0.002946 | |
| std | 0.038535 | 0.037717 | 0.035903 | 0.035522 | 0.035266 | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 50% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 75% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| max | 1.000000 | 1.000000 | 1.000000 | 0.996053 | 1.000000 | |

| | |
|-------|--------------|
| | 187 |
| count | 21892.000000 |
| mean | 0.473689 |
| std | 1.143447 |
| min | 0.000000 |
| 25% | 0.000000 |
| 50% | 0.000000 |
| 75% | 0.000000 |
| max | 4.000000 |

[8 rows x 188 columns]

```
[ ]: df = pd.read_csv('sample_data/mitbih_train.csv', header=None)
df.isnull().sum()
print(df.describe())
```

| | | | | | | |
|-------|--------------|--------------|--------------|--------------|--------------|-----|
| | | 0 | 1 | 2 | 3 | 4 \ |
| count | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | |
| mean | 0.890360 | 0.758160 | 0.423972 | 0.219104 | 0.201127 | |
| std | 0.240909 | 0.221813 | 0.227305 | 0.206878 | 0.177058 | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | 0.921922 | 0.682486 | 0.250969 | 0.048458 | 0.082329 | |
| 50% | 0.991342 | 0.826013 | 0.429472 | 0.166000 | 0.147878 | |
| 75% | 1.000000 | 0.910506 | 0.578767 | 0.341727 | 0.258993 | |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | |

| | | | | | | |
|-------|--------------|--------------|--------------|--------------|--------------|-----|
| | | 5 | 6 | 7 | 8 | 9 \ |
| count | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | |
| mean | 0.210399 | 0.205808 | 0.201773 | 0.198691 | 0.196757 | |

| | | | | | |
|-----|----------|----------|----------|----------|----------|
| std | 0.171909 | 0.178481 | 0.177240 | 0.171778 | 0.168357 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.088416 | 0.073333 | 0.066116 | 0.065000 | 0.068639 |
| 50% | 0.158798 | 0.145324 | 0.144424 | 0.150000 | 0.148734 |
| 75% | 0.287628 | 0.298237 | 0.295391 | 0.290832 | 0.283636 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |

| | | | | | | |
|-------|-----|--------------|--------------|--------------|--------------|---|
| | ... | 178 | 179 | 180 | 181 | \ |
| count | ... | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | |
| mean | ... | 0.005025 | 0.004628 | 0.004291 | 0.003945 | |
| std | ... | 0.044154 | 0.042089 | 0.040525 | 0.038651 | |
| min | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 50% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 75% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| max | ... | 1.000000 | 1.000000 | 1.000000 | 1.000000 | |

| | | | | | | | |
|-------|--------------|--------------|--------------|--------------|--------------|--------------|---|
| | | 182 | 183 | 184 | 185 | 186 | \ |
| count | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | 87554.000000 | |
| mean | | 0.003681 | 0.003471 | 0.003221 | 0.002945 | 0.002807 | |
| std | | 0.037193 | 0.036255 | 0.034789 | 0.032865 | 0.031924 | |
| min | | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 25% | | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 50% | | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| 75% | | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | |
| max | | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 | |

| | |
|-------|--------------|
| | 187 |
| count | 87554.000000 |
| mean | 0.473376 |
| std | 1.143184 |
| min | 0.000000 |
| 25% | 0.000000 |
| 50% | 0.000000 |
| 75% | 0.000000 |
| max | 4.000000 |

[8 rows x 188 columns]

```
[ ]: df = pd.read_csv('sample_data/ptbdb_normal.csv', header=None)

df.isnull().sum()
print(df.describe())
```

| | | | | | | |
|-------|-------------|-------------|-------------|-------------|-------------|---|
| | 0 | 1 | 2 | 3 | 4 | \ |
| count | 4046.000000 | 4046.000000 | 4046.000000 | 4046.000000 | 4046.000000 | |
| mean | 0.979670 | 0.711486 | 0.311677 | 0.119575 | 0.088608 | |
| std | 0.029061 | 0.186376 | 0.183457 | 0.110457 | 0.075760 | |

| | | | | | |
|-----|----------|----------|----------|----------|----------|
| min | 0.782178 | 0.121784 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.964468 | 0.580160 | 0.164112 | 0.028731 | 0.027735 |
| 50% | 1.000000 | 0.726449 | 0.303266 | 0.092655 | 0.079476 |
| 75% | 1.000000 | 0.863699 | 0.436091 | 0.187527 | 0.134861 |
| max | 1.000000 | 1.000000 | 0.985955 | 0.910798 | 0.846591 |

| | | | | | | | |
|-------|-------------|-------------|-------------|-------------|-------------|-----|---|
| | 5 | 6 | 7 | 8 | 9 | ... | \ |
| count | 4046.000000 | 4046.000000 | 4046.000000 | 4046.000000 | 4046.000000 | ... | |
| mean | 0.130843 | 0.159653 | 0.165608 | 0.168005 | 0.170783 | ... | |
| std | 0.081104 | 0.090131 | 0.096094 | 0.102506 | 0.106190 | ... | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | ... | |
| 25% | 0.070776 | 0.094656 | 0.098848 | 0.095594 | 0.094988 | ... | |
| 50% | 0.121144 | 0.134312 | 0.142771 | 0.143002 | 0.145068 | ... | |
| 75% | 0.183746 | 0.216203 | 0.217895 | 0.220572 | 0.223053 | ... | |
| max | 0.770205 | 0.754524 | 0.749095 | 0.729192 | 0.700844 | ... | |

| | | | | | | | |
|-------|-------------|-------------|-------------|-------------|-------------|-----|---|
| | 178 | 179 | 180 | 181 | 182 | ... | \ |
| count | 4046.000000 | 4046.000000 | 4046.000000 | 4046.000000 | 4046.000000 | ... | |
| mean | 0.001540 | 0.001332 | 0.001304 | 0.001220 | 0.000991 | ... | |
| std | 0.018664 | 0.016234 | 0.016668 | 0.016658 | 0.015204 | ... | |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | ... | |
| 25% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | ... | |
| 50% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | ... | |
| 75% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | ... | |
| max | 0.415879 | 0.361283 | 0.383522 | 0.407025 | 0.446281 | ... | |

| | | | | | |
|-------|-------------|-------------|-------------|--------|--------|
| | 183 | 184 | 185 | 186 | 187 |
| count | 4046.000000 | 4046.000000 | 4046.000000 | 4046.0 | 4046.0 |
| mean | 0.000894 | 0.000454 | 0.000474 | 0.0 | 0.0 |
| std | 0.015311 | 0.010834 | 0.011202 | 0.0 | 0.0 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 |
| 25% | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 |
| 50% | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 |
| 75% | 0.000000 | 0.000000 | 0.000000 | 0.0 | 0.0 |
| max | 0.483471 | 0.371502 | 0.376668 | 0.0 | 0.0 |

[8 rows x 188 columns]

```
[ ]: df = pd.read_csv('sample_data/ptbdb_abnormal.csv', header=None)
```

```
df.isnull().sum()
print(df.describe())
```

| | | | | | | | |
|-------|--------------|--------------|--------------|--------------|--------------|-----|---|
| | 0 | 1 | 2 | 3 | 4 | ... | \ |
| count | 10506.000000 | 10506.000000 | 10506.000000 | 10506.000000 | 10506.000000 | ... | |
| mean | 0.975468 | 0.725582 | 0.438306 | 0.290384 | 0.252897 | ... | |
| std | 0.036354 | 0.199030 | 0.262699 | 0.270977 | 0.237004 | ... | |

| | | | | | |
|-----|----------|----------|----------|----------|----------|
| min | 0.624227 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.957325 | 0.586602 | 0.236455 | 0.074064 | 0.088487 |
| 50% | 1.000000 | 0.745646 | 0.404297 | 0.212845 | 0.173046 |
| 75% | 1.000000 | 0.890043 | 0.620889 | 0.427811 | 0.343089 |
| max | 1.000000 | 1.000000 | 1.000000 | 1.000000 | 1.000000 |

| | | | | | |
|-------|--------------|--------------|--------------|--------------|--------------|
| | 5 | 6 | 7 | 8 | 9 \ |
| count | 10506.000000 | 10506.000000 | 10506.000000 | 10506.000000 | 10506.000000 |
| mean | 0.249423 | 0.245668 | 0.247160 | 0.250203 | 0.252396 |
| std | 0.211751 | 0.200159 | 0.194840 | 0.193160 | 0.192341 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | 0.101422 | 0.097175 | 0.096331 | 0.098527 | 0.098041 |
| 50% | 0.180927 | 0.182782 | 0.191793 | 0.197088 | 0.200919 |
| 75% | 0.327582 | 0.328089 | 0.347707 | 0.359378 | 0.370106 |
| max | 1.000000 | 1.000000 | 0.985523 | 0.993213 | 0.997738 |

| | | | | | |
|-------|-----|--------------|--------------|--------------|--------------|
| | ... | 178 | 179 | 180 | 181 \ |
| count | ... | 10506.000000 | 10506.000000 | 10506.000000 | 10506.000000 |
| mean | ... | 0.001055 | 0.001057 | 0.000744 | 0.000554 |
| std | ... | 0.022312 | 0.022585 | 0.017557 | 0.013781 |
| min | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 25% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 50% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| 75% | ... | 0.000000 | 0.000000 | 0.000000 | 0.000000 |
| max | ... | 0.791899 | 0.773743 | 0.789804 | 0.628177 |

| | | | | | |
|-------|--------------|--------------|--------------|--------------|---------|
| | 182 | 183 | 184 | 185 | 186 \ |
| count | 10506.000000 | 10506.000000 | 10506.000000 | 10506.000000 | 10506.0 |
| mean | 0.000533 | 0.000313 | 0.000070 | 0.000074 | 0.0 |
| std | 0.013553 | 0.010901 | 0.003754 | 0.004044 | 0.0 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.0 |
| 25% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.0 |
| 50% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.0 |
| 75% | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.0 |
| max | 0.602033 | 0.644880 | 0.265025 | 0.279310 | 0.0 |

| | |
|-------|---------|
| | 187 |
| count | 10506.0 |
| mean | 1.0 |
| std | 0.0 |
| min | 1.0 |
| 25% | 1.0 |
| 50% | 1.0 |
| 75% | 1.0 |
| max | 1.0 |

[8 rows x 188 columns]

5 Prueba del primero modelo

5.1 Regresión logística sin balanceo

```
[ ]: df = pd.read_csv('sample_data/mitbih_train.csv', header=None)

X = df.iloc[:, :187]
y = df[187]
```

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
    ↪random_state=42)

log_reg = LogisticRegression(max_iter=1000, random_state=42)

cv_scores = cross_val_score(log_reg, X_train, y_train, cv=5, scoring='accuracy')

log_reg.fit(X_train, y_train)

y_pred = log_reg.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)
print(f'Precisión del modelo: {accuracy}')

print(f'Precisión promedio con validación cruzada: {cv_scores.mean()}')
print(f'Precisión en cada partición de la validación cruzada: {cv_scores}')
```

Precisión del modelo: 0.9139604827349906

Precisión promedio con validación cruzada: 0.9131952967948511

Precisión en cada partición de la validación cruzada: [0.91164953 0.91458639
0.91147915 0.91539528 0.91286612]

```
[ ]: report = classification_report(y_test, y_pred)
print("Reporte de clasificación:")
print(report)
```

Reporte de clasificación:

| | precision | recall | f1-score | support |
|-----|-----------|--------|----------|---------|
| 0.0 | 0.92 | 0.98 | 0.95 | 21828 |
| 1.0 | 0.85 | 0.39 | 0.53 | 627 |
| 2.0 | 0.64 | 0.33 | 0.43 | 1704 |
| 3.0 | 0.63 | 0.20 | 0.30 | 200 |

| | | | | | |
|--------------|------|------|------|------|-------|
| | 4.0 | 0.95 | 0.88 | 0.92 | 1908 |
| accuracy | | | | 0.91 | 26267 |
| macro avg | 0.80 | 0.56 | 0.63 | | 26267 |
| weighted avg | 0.90 | 0.91 | 0.90 | | 26267 |


```
[ ]: conf_matrix = confusion_matrix(y_test, y_pred)

print("Matriz de confusión:")
print(conf_matrix)

plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Reds', xticklabels=["N", "S", "V", "F", "Q"], yticklabels=["N", "S", "V", "F", "Q"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

Matriz de confusión:

```
[[21486   33   252    4   53]
 [  355   243    26    1    2]
 [ 1093    11   554   18   28]
 [  146     0    14   40    0]
 [  202     0    22    0 1684]]
```



6 Modelos

6.1 7 modelos balanceados por RandomUnderSample

```
[ ]: df = pd.read_csv('sample_data/mitbih_train.csv', header=None)

X = df.iloc[:, :187].values
y = df[187].values

n_folds = 5
kf = KFold(n_splits=n_folds, shuffle=True, random_state=42)

models = {
    'Random Forest': RandomForestClassifier(random_state=42),
    'SVM': SVC(random_state=42),
    'K-Nearest Neighbors': KNeighborsClassifier(),
    'Decision Tree': DecisionTreeClassifier(random_state=42),
    'Naive Bayes': GaussianNB(),
```



```

'Gradient Boosting': GradientBoostingClassifier(random_state=42),
'AdaBoost': AdaBoostClassifier(random_state=42)
}

results = {}

for model_name, model in models.items():
    print(f"Evaluando modelo: {model_name}")

    accuracy_total = 0

    for train_index, test_index in kf.split(X):
        X_train, X_test = X[train_index], X[test_index]
        y_train, y_test = y[train_index], y[test_index]

        undersample = RandomUnderSampler(sampling_strategy='auto',
↪random_state=42)
        X_res, y_res = undersample.fit_resample(X_train, y_train)

        model.fit(X_res, y_res)

        y_pred = model.predict(X_test)
        accuracy = accuracy_score(y_test, y_pred)
        accuracy_total += accuracy

    accuracy_avg = accuracy_total / n_folds
    results[model_name] = accuracy_avg
    print(f"Precisión promedio para {model_name}: {accuracy_avg:.4f}\n")

    report = classification_report(y_test, y_pred)
    print("Reporte de clasificación:")
    print(report)

    conf_matrix = confusion_matrix(y_test, y_pred)

    plt.figure(figsize=(10, 7))
    sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Reds',
↪xticklabels=["N", "S", "V", "F", "Q"], yticklabels=["N", "S", "V", "F", "Q"])
    plt.xlabel('Predicted')
```

```

plt.ylabel('Actual')
plt.title(f'Confusion Matrix for {model_name}')
plt.show()


best_model_name = max(results, key=results.get)
print(f"\nEl modelo más efectivo es: {best_model_name} con una precisión_
↳promedio de {results[best_model_name]:.4f}")

```

Evaluando modelo: Random Forest
Precisión promedio para Random Forest: 0.8511

Reporte de clasificación:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.99 | 0.85 | 0.91 | 14452 |
| 1.0 | 0.34 | 0.80 | 0.47 | 465 |
| 2.0 | 0.62 | 0.92 | 0.74 | 1183 |
| 3.0 | 0.13 | 0.88 | 0.22 | 127 |
| 4.0 | 0.91 | 0.95 | 0.93 | 1283 |
| accuracy | | | 0.86 | 17510 |
| macro avg | 0.59 | 0.88 | 0.65 | 17510 |
| weighted avg | 0.93 | 0.86 | 0.88 | 17510 |




Evaluando modelo: SVM

Precisión promedio para SVM: 0.8082

Reporte de clasificación:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.98 | 0.80 | 0.88 | 14452 |
| 1.0 | 0.33 | 0.75 | 0.46 | 465 |
| 2.0 | 0.59 | 0.86 | 0.70 | 1183 |
| 3.0 | 0.08 | 0.87 | 0.14 | 127 |
| 4.0 | 0.83 | 0.91 | 0.87 | 1283 |
| accuracy | | | 0.81 | 17510 |
| macro avg | 0.56 | 0.84 | 0.61 | 17510 |
| weighted avg | 0.92 | 0.81 | 0.85 | 17510 |




Evaluando modelo: K-Nearest Neighbors

Precisión promedio para K-Nearest Neighbors: 0.7586

Reporte de clasificación:

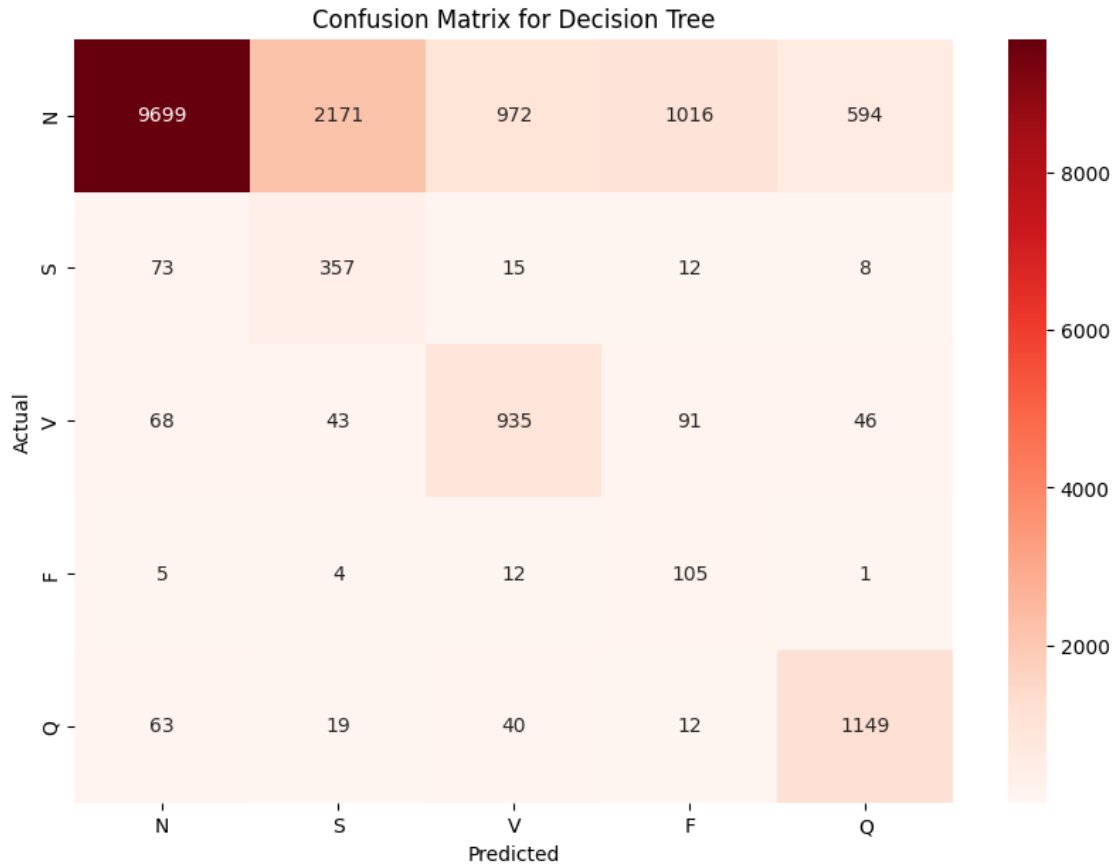
| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.98 | 0.74 | 0.84 | 14452 |
| 1.0 | 0.19 | 0.81 | 0.31 | 465 |
| 2.0 | 0.51 | 0.84 | 0.63 | 1183 |
| 3.0 | 0.09 | 0.88 | 0.16 | 127 |
| 4.0 | 0.77 | 0.93 | 0.85 | 1283 |
| accuracy | | | 0.76 | 17510 |
| macro avg | 0.51 | 0.84 | 0.56 | 17510 |
| weighted avg | 0.91 | 0.76 | 0.81 | 17510 |



Evaluando modelo: Decision Tree
 Precisión promedio para Decision Tree: 0.7101

Reporte de clasificación:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.98 | 0.67 | 0.80 | 14452 |
| 1.0 | 0.14 | 0.77 | 0.23 | 465 |
| 2.0 | 0.47 | 0.79 | 0.59 | 1183 |
| 3.0 | 0.08 | 0.83 | 0.15 | 127 |
| 4.0 | 0.64 | 0.90 | 0.75 | 1283 |
| accuracy | | | 0.70 | 17510 |
| macro avg | 0.46 | 0.79 | 0.50 | 17510 |
| weighted avg | 0.89 | 0.70 | 0.76 | 17510 |



Evaluando modelo: Naive Bayes

Precisión promedio para Naive Bayes: 0.1792

Reporte de clasificación:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.92 | 0.12 | 0.22 | 14452 |
| 1.0 | 0.52 | 0.08 | 0.14 | 465 |
| 2.0 | 0.22 | 0.19 | 0.20 | 1183 |
| 3.0 | 0.00 | 0.00 | 0.00 | 127 |
| 4.0 | 0.09 | 1.00 | 0.16 | 1283 |
| accuracy | | | 0.19 | 17510 |
| macro avg | 0.35 | 0.28 | 0.14 | 17510 |
| weighted avg | 0.80 | 0.19 | 0.21 | 17510 |


/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471:
 UndefinedMetricWarning: Precision and F-score are ill-defined and being set to
 0.0 in labels with no predicted samples. Use `zero_division` parameter to

control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))  
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))  
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471:  
UndefinedMetricWarning: Precision and F-score are ill-defined and being set to  
0.0 in labels with no predicted samples. Use `zero_division` parameter to  
control this behavior.
```

```
_warn_prf(average, modifier, msg_start, len(result))
```




Evaluando modelo: Gradient Boosting

Precisión promedio para Gradient Boosting: 0.7945

Reporte de clasificación:

| | precision | recall | f1-score | support | |
|--|-----------|--------|----------|---------|-------|
| | 0.0 | 0.98 | 0.78 | 0.87 | 14452 |

| | | | | |
|--------------|------|------|------|-------|
| 1.0 | 0.24 | 0.80 | 0.37 | 465 |
| 2.0 | 0.56 | 0.89 | 0.69 | 1183 |
| 3.0 | 0.09 | 0.87 | 0.17 | 127 |
| 4.0 | 0.85 | 0.93 | 0.89 | 1283 |
| accuracy | | | 0.80 | 17510 |
| macro avg | 0.55 | 0.86 | 0.60 | 17510 |
| weighted avg | 0.92 | 0.80 | 0.84 | 17510 |



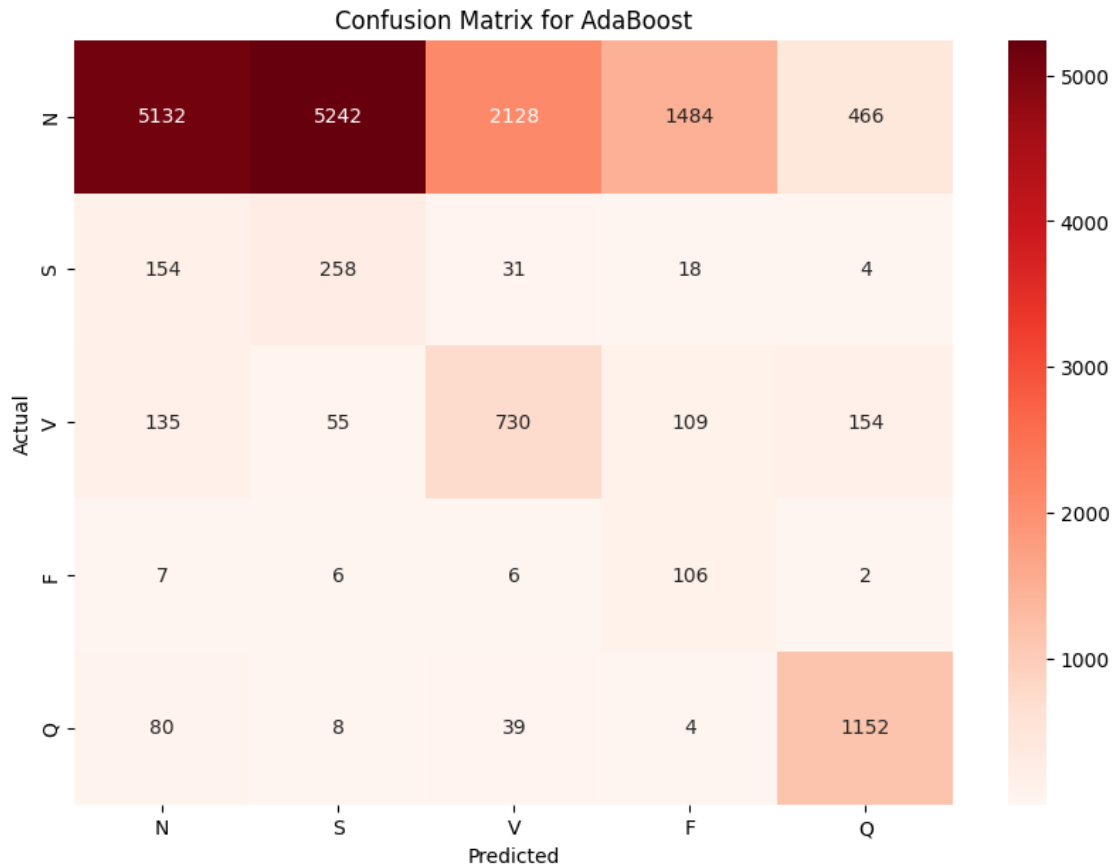
Evaluando modelo: AdaBoost

Precisión promedio para AdaBoost: 0.4800

Reporte de clasificación:

| | precision | recall | f1-score | support |
|-----|-----------|--------|----------|---------|
| 0.0 | 0.93 | 0.36 | 0.51 | 14452 |
| 1.0 | 0.05 | 0.55 | 0.09 | 465 |
| 2.0 | 0.25 | 0.62 | 0.35 | 1183 |
| 3.0 | 0.06 | 0.83 | 0.11 | 127 |

| | | | | | |
|--------------|------|------|------|------|-------|
| | 4.0 | 0.65 | 0.90 | 0.75 | 1283 |
| accuracy | | | | 0.42 | 17510 |
| macro avg | 0.39 | 0.65 | 0.36 | | 17510 |
| weighted avg | 0.83 | 0.42 | 0.51 | | 17510 |



El modelo más efectivo es: Random Forest con una precisión promedio de 0.8511

6.2 SMOTE ramdonForest

```
[ ]: from imblearn.over_sampling import SMOTE

df = pd.read_csv('sample_data/mitbih_train.csv', header=None)

X = df.iloc[:, :187].values
y = df[187].values

n_folds = 5
```

```

kf = KFold(n_splits=n_folds, shuffle=True, random_state=42)

model_name = 'Random Forest'
model = RandomForestClassifier(random_state=42)

accuracy_total = 0

for train_index, test_index in kf.split(X):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]

    smote = SMOTE(random_state=42)
    X_res, y_res = smote.fit_resample(X_train, y_train)

    model.fit(X_res, y_res)

    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    accuracy_total += accuracy

accuracy_avg = accuracy_total / n_folds
print(f"Precisión promedio para {model_name}: {accuracy_avg:.4f}\n")

report = classification_report(y_test, y_pred)
print("Reporte de clasificación:")
print(report)

conf_matrix = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Reds', xticklabels=["N", "S", "V", "F", "Q"], yticklabels=["N", "S", "V", "F", "Q"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix for {model_name}')
plt.show()

```

7 Solución

7.1 Modelo 1 = Normal + Abnormal

```
[ ]: from sklearn.model_selection import StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, classification_report, \
    accuracy_score
from imblearn.under_sampling import RandomUnderSampler
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

[ ]: df_normal = pd.read_csv('sample_data/ptbdb_normal.csv', header=None)
df_abnormal = pd.read_csv('sample_data/ptbdb_abnormal.csv', header=None)

df_normal[187] = 0
df_abnormal[187] = 1

df_combined = pd.concat([df_normal, df_abnormal], axis=0).reset_index(drop=True)

X = df_combined.iloc[:, :187]
y = df_combined[187]
```

7.1.1 Balanceo mediante observaciones normales de mitbih_train

```
[ ]: df_mitbih = pd.read_csv('sample_data/mitbih_train.csv', header=None)

df_mitbih_normal = df_mitbih[df_mitbih[187] == 0]

normales_faltantes = 2000

df_mitbih_normal_extra = df_mitbih_normal.sample(n=normales_faltantes, \
    random_state=42)

df_combined_balanced = pd.concat([df_combined, df_mitbih_normal_extra], axis=0). \
    reset_index(drop=True)

X_balanced = df_combined_balanced.iloc[:, :187]
y_balanced = df_combined_balanced[187]

class_counts_balanced = y_balanced.value_counts()
print("Número de ejemplos por clase después de balancear:")
for cls, count in class_counts_balanced.items():
    print(f"Clase {cls}: {count} ejemplos")
```

Número de ejemplos por clase después de balancear:
Clase 1.0: 10506 ejemplos

Clase 0.0: 6046 ejemplos

7.1.2 Entrenamiento

```
[ ]: X_balanced = df_combined_balanced.iloc[:, :187]
y_balanced = df_combined_balanced[187]

under_sampler = RandomUnderSampler(random_state=42)
random_forest_model = RandomForestClassifier(random_state=42)

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

y_true = []
y_pred = []

for train_index, test_index in cv.split(X_balanced, y_balanced):
    X_train, X_test = X_balanced.iloc[train_index], X_balanced.iloc[test_index]
    y_train, y_test = y_balanced.iloc[train_index], y_balanced.iloc[test_index]

    X_train_under, y_train_under = under_sampler.fit_resample(X_train, y_train)

    random_forest_model.fit(X_train_under, y_train_under)

    y_pred_fold = random_forest_model.predict(X_test)

    y_true.extend(y_test)
    y_pred.extend(y_pred_fold)

y_true = pd.Series(y_true)
y_pred = pd.Series(y_pred)

conf_matrix_rf = confusion_matrix(y_true, y_pred)
report_rf = classification_report(y_true, y_pred)
accuracy_rf = accuracy_score(y_true, y_pred)

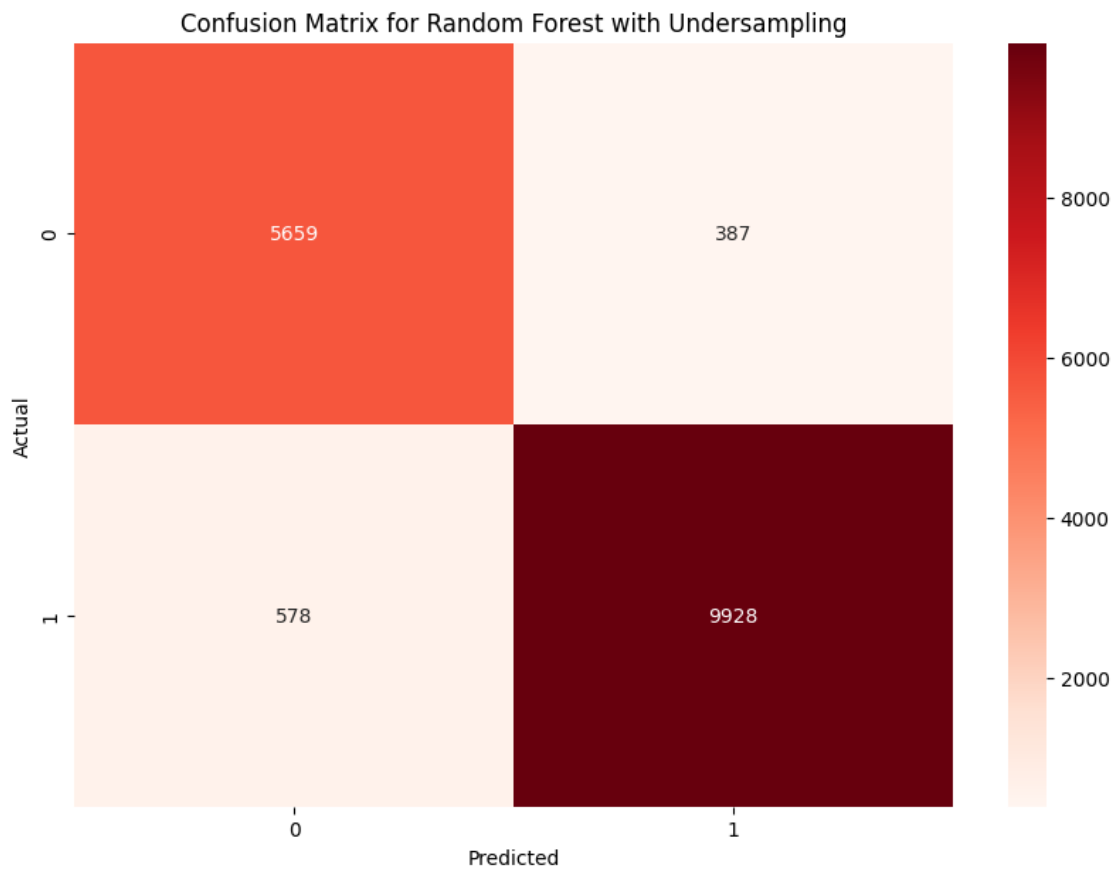
plt.figure(figsize=(10, 7))
```

```

sns.heatmap(conf_matrix_rf, annot=True, fmt='d', cmap='Reds', xticklabels=["0", "1"], yticklabels=["0", "1"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix for Random Forest with Undersampling')
plt.show()

print("\nReporte de Clasificación - Random Forest con Undersampling:")
print(report_rf)
print(f"Exactitud (Accuracy): {accuracy_rf:.2f}")

```



Reporte de Clasificación - Random Forest con Undersampling:

| | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| 0.0 | 0.91 | 0.94 | 0.92 | 6046 |
| 1.0 | 0.96 | 0.94 | 0.95 | 10506 |
| accuracy | | | 0.94 | 16552 |

| | | | | |
|--------------|------|------|------|-------|
| macro avg | 0.93 | 0.94 | 0.94 | 16552 |
| weighted avg | 0.94 | 0.94 | 0.94 | 16552 |

Exactitud (Accuracy): 0.94

7.1.3 Prueba con mitbih_test

```
[ ]: df_test = pd.read_csv('sample_data/mitbih_test.csv', header=None)

X_test = df_test.iloc[:, :187]
y_test = df_test[187]

y_test = y_test.replace({2: 1, 3: 1, 4: 1})


random_forest_model = RandomForestClassifier()
random_forest_model.fit(X_balanced, y_balanced)

y_pred_test = random_forest_model.predict(X_test)

conf_matrix_test = confusion_matrix(y_test, y_pred_test)
report_test = classification_report(y_test, y_pred_test)
accuracy_test = accuracy_score(y_test, y_pred_test)

plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix_test, annot=True, fmt='d', cmap='Reds',
            xticklabels=["0", "1"], yticklabels=["0", "1"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for Random Forest Classifier')
plt.show()

print("\nReporte de Clasificación mitbih_test - Random Forest:")
print(report_test)
print(f"Exactitud (Accuracy): {accuracy_test:.2f}")
```



Reporte de Clasificación mitbih_test - Random Forest:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0.0 | 0.91 | 0.94 | 0.92 | 18118 |
| 1.0 | 0.66 | 0.52 | 0.59 | 3774 |
| accuracy | | | 0.87 | 21892 |
| macro avg | 0.78 | 0.73 | 0.75 | 21892 |
| weighted avg | 0.86 | 0.87 | 0.87 | 21892 |

Exactitud (Accuracy): 0.87

7.2 Modelo 2

```
[ ]: df = pd.read_csv('sample_data/mitbih_train.csv', header=None)

X = df.iloc[:, :187]
y = df[187]
```

```
[ ]: mask = y != 0
      X_filtered = X[mask]
      y_filtered = y[mask]

      df_abnormal = df[mask]

      print(X_filtered.shape)
      print(y_filtered.shape)
```

```
(15083, 187)
(15083,)
```

```
[ ]: X = df_abnormal.iloc[:, :187]
      y = df_abnormal[187]
```

7.2.1 Entrenamiento

```
[ ]: from sklearn.model_selection import cross_val_predict
```


```
[ ]: random_forest_model = RandomForestClassifier()

      y_pred_rf = cross_val_predict(random_forest_model, X, y, cv=5)

      conf_matrix_rf = confusion_matrix(y, y_pred_rf)

      report_rf = classification_report(y, y_pred_rf)
      accuracy_rf = accuracy_score(y, y_pred_rf)

      plt.figure(figsize=(10, 7))
      sns.heatmap(conf_matrix_rf, annot=True, fmt='d', cmap='Reds', xticklabels=[
          ↪ "S", "Q", "V", "F"], yticklabels=["S", "Q", "V", "F"])
      plt.xlabel('Predicted')
      plt.ylabel('Actual')
      plt.title(f'Confusion Matrix for {random_forest_model}')
      plt.show()
      print("\nReporte de Clasificación - Regresión Logística:")
      print(report_rf)
      print(f"Exactitud (Accuracy): {accuracy_rf:.2f}")
```

Reporte de Clasificación - Regresión Logística:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1.0 | 0.97 | 0.94 | 0.95 | 2223 |
| 2.0 | 0.96 | 0.98 | 0.97 | 5788 |
| 3.0 | 0.92 | 0.82 | 0.87 | 641 |
| 4.0 | 0.99 | 0.99 | 0.99 | 6431 |
| accuracy | | | 0.97 | 15083 |
| macro avg | 0.96 | 0.93 | 0.95 | 15083 |
| weighted avg | 0.97 | 0.97 | 0.97 | 15083 |

Exactitud (Accuracy): 0.97

7.2.2 Prueba con mitbih_test

```
[ ]: df_test = pd.read_csv('sample_data/mitbih_test.csv', header=None)

X_test = df_test.iloc[:, :187]
y_test = df_test[187]

mask_test = y_test != 0
X_test_filtered = X_test[mask_test]
y_test_filtered = y_test[mask_test]


random_forest_model.fit(X, y)

y_pred_test = random_forest_model.predict(X_test_filtered)

conf_matrix_test = confusion_matrix(y_test_filtered, y_pred_test)
report_test = classification_report(y_test_filtered, y_pred_test)
accuracy_test = accuracy_score(y_test_filtered, y_pred_test)

plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix_test, annot=True, fmt='d', cmap='Reds',
            xticklabels=["S", "Q", "V", "F"], yticklabels=["S", "Q", "V", "F"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix for Random Forest Classifier')
plt.show()

print("\nReporte de Clasificación - Random Forest:")
print(report_test)
print(f"Exactitud (Accuracy): {accuracy_test:.2f}")
```



Reporte de Clasificación - Random Forest:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1.0 | 0.98 | 0.92 | 0.95 | 556 |
| 2.0 | 0.95 | 0.98 | 0.96 | 1448 |
| 3.0 | 0.84 | 0.81 | 0.82 | 162 |
| 4.0 | 0.99 | 0.99 | 0.99 | 1608 |
| accuracy | | | 0.97 | 3774 |
| macro avg | 0.94 | 0.93 | 0.93 | 3774 |
| weighted avg | 0.97 | 0.97 | 0.97 | 3774 |

Exactitud (Accuracy): 0.97

8 Ramdom forest hiperparametros

```
[ ]: df_mitbih = pd.read_csv('sample_data/mitbih_train.csv', header=None)

df_mitbih_normal = df_mitbih[df_mitbih[187] == 0]

normales_faltantes = 6460

df_mitbih_normal_extra = df_mitbih_normal.sample(n=normales_faltantes,
↪random_state=42)

df_combined_balanced = pd.concat([df_combined, df_mitbih_normal_extra], axis=0).
↪reset_index(drop=True)

X_balanced = df_combined_balanced.iloc[:, :187]
y_balanced = df_combined_balanced[187]

from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [10, 20],
    'min_samples_split': [2, 5],
    'min_samples_leaf': [1, 2],
    'bootstrap': [True, False]
}

grid_search = GridSearchCV(estimator=random_forest_model, param_grid=param_grid,
↪cv=3, n_jobs=-1, verbose=1, scoring='accuracy')
grid_search.fit(X_balanced, y_balanced)

best_params = grid_search.best_params_
best_model = grid_search.best_estimator_
print(f"Mejores parámetros: {best_params}")
```

Fitting 3 folds for each of 32 candidates, totalling 96 fits

Mejores parámetros: {'bootstrap': False, 'max_depth': 20, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 100}

```
[ ]: y_pred_best_rf_random = cross_val_predict(best_model, X_balanced, y_balanced,
↪cv=5)

conf_matrix_best_rf_random = confusion_matrix(y_balanced, y_pred_best_rf_random)
report_best_rf_random = classification_report(y_balanced, y_pred_best_rf_random)
accuracy_best_rf_random = accuracy_score(y_balanced, y_pred_best_rf_random)
```

```

plt.figure(figsize=(10, 7))
sns.heatmap(conf_matrix_best_rf_random, annot=True, fmt='d', cmap='Reds',
            xticklabels=["N", "AB"], yticklabels=["N", "AB"])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title(f'Confusion Matrix for {best_model}')
plt.show()
print("\nReporte de Clasificación - Random Forest con Mejor Parámetro
      (Randomized Search):")
print(report_best_rf_random)
print(f"Exactitud (Accuracy): {accuracy_best_rf_random:.2f}")

```

8.1 mas hiperparametros

```

[ ]: import pandas as pd

df_normal = pd.read_csv('sample_data/ptbdb_normal.csv', header=None)
df_abnormal = pd.read_csv('sample_data/ptbdb_abnormal.csv', header=None)

df_normal[187] = 0
df_abnormal[187] = 1

df_combined = pd.concat([df_normal, df_abnormal], axis=0).reset_index(drop=True)

X = df_combined.iloc[:, :187]
y = df_combined[187]

class_counts = y.value_counts()

print("Número de ejemplos por clase:")
for cls, count in class_counts.items():
    print(f"Clase {cls}: {count} ejemplos")

```

Número de ejemplos por clase:
 Clase 1: 10506 ejemplos
 Clase 0: 4046 ejemplos

```

[ ]: df_mitbih = pd.read_csv('sample_data/mitbih_train.csv', header=None)

df_mitbih_normal = df_mitbih[df_mitbih[187] == 0]

normales_faltantes = 6460

```

```

df_mitbih_normal_extra = df_mitbih_normal.sample(n=normales_faltantes,
↳random_state=42)

df_combined_balanced = pd.concat([df_combined, df_mitbih_normal_extra], axis=0).
↳reset_index(drop=True)

X_balanced = df_combined_balanced.iloc[:, :187]
y_balanced = df_combined_balanced[187]

```

8.1.1 No corrio, mucho costo computacional requerido

```

[ ]: from sklearn.model_selection import GridSearchCV

param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, 30, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'bootstrap': [True, False],
    'class_weight': ['balanced', {0: 1, 1: 4}]
}

rf_model = RandomForestClassifier()

grid_search = GridSearchCV(estimator=rf_model, param_grid=param_grid, cv=5,
↳scoring='accuracy', n_jobs=-1, verbose=2)
grid_search.fit(X_balanced, y_balanced)

print(f"Mejores parámetros encontrados: {grid_search.best_params_}")

best_model = grid_search.best_estimator_
y_pred_test_best = best_model.predict(X_test)

conf_matrix_test_best = confusion_matrix(y_test, y_pred_test_best)
report_test_best = classification_report(y_test, y_pred_test_best)
accuracy_test_best = accuracy_score(y_test, y_pred_test_best)

print("\nReporte de Clasificación con el Mejor Modelo:")
print(report_test_best)

```

```
print(f"Exactitud (Accuracy) con el Mejor Modelo: {accuracy_test_best:.2f}")
```

Fitting 5 folds for each of 432 candidates, totalling 2160 fits